

DS 598

Introduction to RL

Xuezhou Zhang

Chapter 5: Policy-based RL (continued)

The REINFORCE algorithm

1. Initialize θ_0
2. For iteration $t = 0, \dots, T$
 - 1) Run π_{θ_t} and collect trajectories τ_1, \dots, τ_n
 - 2) Estimate the PG by

$$g_t = \frac{1}{n} \sum_{i=1}^n \left[\sum_{h=0}^{\infty} \nabla_{\theta} \log \pi(a_{i;h} | s_{i;h}) R(\tau_i) \right]$$

- 3) Do SGD update $\theta_{t+1} = \theta_t + \alpha_t g_t$

The REINFORCE algorithm

- Pros:

- ✓ Convergence
- ✓ Conceptually simple

- Cons:

- ❖ Only works with stochastic policies
- ❖ On-policy -> Sample inefficient
- ❖ High Variance $\mathbb{E} \left[\left\| g_t - \nabla_{\theta} J(\pi_{\theta}) \right\|_2^2 \right]$

High Variance

- $\mathbb{E} \left[\|g_t - \nabla_{\theta} J(\pi_{\theta})\|_2^2 \right]$ can be up to $(1 - \gamma)^3$.

$$g_t = \frac{1}{n} \sum_{i=1}^n \left[\sum_{h=0}^{\infty} \nabla_{\theta} \log \pi(a_{i,h} | s_{i,h}) R(\tau_i) \right]$$

Reducing Variance

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{h=0}^{\infty} \nabla_{\theta} \log \pi(a_h | s_h) R(\tau) \right]$$

- Key Observation: $\mathbb{E}_{a \sim \pi(s)} [\nabla_{\theta} \log \pi(a | s) f(s)] = 0$

- Therefore, $\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{h=0}^{\infty} \nabla_{\theta} \log \pi(a_h | s_h) \sum_{i=0}^{\infty} \gamma^h r(s_i, a_i) \right]$

Reducing Variance

- A “good” baseline: $b(s) = V^{\pi_\theta}(s)$.

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{h=0}^{\infty} \nabla_\theta \log \pi(a_h | s_h) (Q^{\pi_\theta}(s_h, a_h) - b(s_h)) \right]$$

Advantage function: $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$

Reducing Variance

- Sure, but how do we estimate $A^{\pi_\theta}(s, a)$?
- Estimate $V^{\pi_\theta}(s)$ using TD-learning.

$$\hat{V}^\pi = \arg \min_f \sum_{i=1}^N \sum_{h=1}^H [f(s_{i;h}) - (r_{i;h} + \gamma r_{i;h+1} + \gamma^2 r_{i;h+2} + \dots)]^2$$

- Then, $A^{\pi_\theta}(s_h, a_h) \approx r_{i;h} + \gamma r_{i;h+1} + \gamma^2 r_{i;h+2} + \dots - \hat{V}^\pi(s_{i;h})$
- This is the so-called Advantage Actor-Critic (A2C) algorithm.

The A2C algorithm

1. Initialize θ_0

2. For iteration $t = 0, \dots, T$

1) Run π_{θ_t} and collect trajectories τ_1, \dots, τ_n

2) Update Critic: $\hat{V} = \arg \min_f \sum_{i=1}^N \sum_{h=1}^H [f(s_{i;h}) - (r_{i;h} + \gamma r_{i;h+1} + \gamma^2 r_{i;h+2} + \dots)]^2$

3) Estimate the PG by

$$g_t = \frac{1}{n} \sum_{i=1}^n \left[\sum_{h=0}^{\infty} \nabla_{\theta} \log \pi(a_{i;h} | s_{i;h}) (R(s_{i;h}, a_{i;h}) - \hat{V}(s_{i;h})) \right]$$

4) Do SGD update $\theta_{t+1} = \theta_t + \alpha_t g_t$

The A2C algorithm

- Pros:

- ✓ Convergence
- ✓ Conceptually simple
- ✓ Low PG Variance

- Cons:

- ❖ Only works with stochastic policies
- ❖ On-policy -> Sample inefficient

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{h=0}^{\infty} \nabla_{\theta} \log \pi(a_h | s_h) R(\tau) \right]$$

Can we make use of off-policy data?

importance sampling

$$\begin{aligned} E_{x \sim p(x)}[f(x)] &= \int p(x) f(x) dx \\ &= \int \frac{q(x)}{q(x)} p(x) f(x) dx \\ &= \int q(x) \frac{p(x)}{q(x)} f(x) dx \\ &= E_{x \sim q(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \end{aligned}$$

IS-based Off-Policy Policy Gradient

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{h=0}^{\infty} \nabla_{\theta} \log \pi_{\theta}(a_h | s_h) R(\tau) \right]$$

Exponentially large -> blows up the variance

The IS-PG algorithm

- Pros:

- ✓ Convergence
- ✓ Conceptually simple
- ✓ Low PG Variance
- ✓ Can be off-policy

- Cons:

- ❖ Only works with stochastic policies

Deterministic Policy Gradient Theorem

- Deterministic Policy: $a = \pi_{\theta}(s)$
- Deterministic Policy Gradient Theorem

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{s \sim d^{\pi_{\theta}}} [\nabla_{\theta} \pi(s) \nabla_a Q^{\pi_{\theta}}(s, \pi_{\theta}(a))]$$

Deterministic Policy Gradient Theorem

- How to learn Q^{π_θ} ?
- For stochastic policies: $Q^\pi(s_t, a_t) = \mathbb{E}_{r_t, s_{t+1} \sim E} [r(s_t, a_t) + \gamma \mathbb{E}_{a_{t+1} \sim \pi} [Q^\pi(s_{t+1}, a_{t+1})]]$
- For deterministic policies: $Q^\mu(s_t, a_t) = \mathbb{E}_{r_t, s_{t+1} \sim E} [r(s_t, a_t) + \gamma Q^\mu(s_{t+1}, \mu(s_{t+1}))]$
- Can do off-policy learning!

Deep Deterministic Policy Gradient (DDPG)

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q$, $\theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M **do**

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation state s_1

for t = 1, T **do**

 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in R

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

technically wrong

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

end for
end for
